

# Parallel and Distributed Computing for Optimization and Statistics

Chris McCord<sup>1</sup>

MIT

February 1, 2018

---

<sup>1</sup>Special thanks to Sebastien Martin for providing many of the slides!

# Heavy Computations

In Optimization and Statistics, we often need a lot of computational power:

- Machine Learning on large datasets
- Simulations that require many iterations to converge
- Hard optimization problems, mixed integer programming

# Heavy Computations

In Optimization and Statistics, we often need a lot of computational power:

- Machine Learning on large datasets
- Simulations that require many iterations to converge
- Hard optimization problems, mixed integer programming

Or repetitive computations:

- Parameter tuning
- Benchmarking

# Limitations of a Personal Computer

Using your personal computer may seem simple, but there are serious limitations:

- Limited **memory** (Big Data, large matrices...)

# Limitations of a Personal Computer

Using your personal computer may seem simple, but there are serious limitations:

- Limited **memory** (Big Data, large matrices...)
- Limited **computational power**.

# Limitations of a Personal Computer

Using your personal computer may seem simple, but there are serious limitations:

- Limited **memory** (Big Data, large matrices...)
- Limited **computational power**.
- Limited **number of cores**.

# Limitations of a Personal Computer

Using your personal computer may seem simple, but there are serious limitations:

- Limited **memory** (Big Data, large matrices...)
- Limited **computational power**.
- Limited **number of cores**.
- Limited **time** available. (You, presumably, have better things to do than sitting and watching your program run.)

# Resources

There are several ways to have access to a remote computer or a cluster.

- Another personal computer you own.
- Athena at MIT.
- Resources of your department/lab (Engaging at Sloan).
- Cloud computing as a service (Amazon AWS, Google Cloud Computing...).



# Distributed Computing, Parallel Computing, Remote Computing, etc.

Many terms are used to described this type of computing, and the distinctions between them aren't always clear.

- **Remote computing** - execute commands on a remote machine

# Distributed Computing, Parallel Computing, Remote Computing, etc.

Many terms are used to describe this type of computing, and the distinctions between them aren't always clear.

- **Remote computing** - execute commands on a remote machine
- **Parallel computing** - execute multiple commands simultaneously

# Distributed Computing, Parallel Computing, Remote Computing, etc.

Many terms are used to describe this type of computing, and the distinctions between them aren't always clear.

- **Remote computing** - execute commands on a remote machine
- **Parallel computing** - execute multiple commands simultaneously
- **Distributed computing** - parallel computing without shared memory

# Why Should I Use This?


For research:

- Tackle **bigger data sets** in Stats and Optimization.
- **Parallelize** your computations.
- Longer computational times: **run overnight** or for a whole week!
- Can be very simple to use with interactive sessions (**RStudio**...)

In general:


- **Valuable skill**, used everywhere in industry.

# How does it work: using a remote computer



- We use **SSH** (see lecture 1) to run commands on a remote machine through our computer.
- We can use the **shell** to do almost anything on the remote computer: create files, run a program, use Julia...
- It is also possible to use GUI applications like **R-Studio** or **Matlab**.
- It works with any computer, including your own machines.


# How does it work: using a cluster



When a computing cluster is available, we can run multiple “jobs” thanks to a **job management system**.

- Different job management systems exist, we will use the example of Slurm, which the cluster **Engaging** uses.
- More complex to use, but far more **powerful**.

# How does it work: using a cluster



When a computing cluster is available, we can run multiple “jobs” thanks to a **job management system**.

- Different job management systems exist, we will use the example of Slurm, which the cluster **Engaging** uses.
- More complex to use, but far more **powerful**.

# Engaging

Engaging is a powerful computing cluster for MIT Sloan affiliates, request an account at `stshelp@mit.edu`. It uses the job management system **Slurm**, that we will use in this presentation.

## Connecting to Engaging

**Hostname** `eosloan.mit.edu`

**Username** MIT Kerberos username

**Password** MIT Kerberos password

[wikis.mit.edu/confluence/display/sloanrc/Engaging+Platform](https://wikis.mit.edu/confluence/display/sloanrc/Engaging+Platform)



# Athena

Any MIT affiliate can access the Athena distributed computing environment, you can use it to follow if you do not have access to Engaging.

## Connecting to Athena

**Hostname** athena.dialup.mit.edu

**Username** MIT Kerberos username

**Password** MIT Kerberos password

<http://web.mit.edu/dialup/www/ssh.html>

# SSH

As seen in Lecture 1, we can connect to the remote machine through SSH:

```
ssh <user_name>@eosloan.mit.edu
```


# File Transfer

- Downloading from the Internet (useful to install software, or download datasets ): `wget`
- Transferring files between the local and remote machines: Secure Copy `scp`

```
scp file_name.csv <user_name>@eosloan.mit.edu:
```

# Using Git/GitHub to Synchronize Code

If you need to run code on the remote machine, an elegant way is to use git and GitHub:



## For Advanced Users

Feel free to read up if you're interested!

- **sshfs** allows you to access the files of the remote machine as if they were on your own (magic)!
- **screen** allows you to keep your session active after you disconnect.
- **tmux** (terminal multiplexer) is like screen with lots of additional functionalities (panes...)


# Running Jobs on a Computing Cluster

## Important!

Do not run jobs/processes on the login node!

- When we login to Engaging, we are actually logging into a special **login node**.
- The login nodes are not designed for performing computations, but rather just for submitting jobs to the job management system.
- The actual work will be done by the **computing nodes**.

# Running Jobs on a Computing Cluster




**Job** A program to be executed. (running a R/Julia/Python source file, a bash script, etc..)

**Resources** Resources required to run the job (memory, cpus, time, software requirements...)

The Job/Resource manager allocates the available resources of the computing nodes of the cluster to the different users.

# Running Jobs on a Computing Cluster




**Job** A program to be executed. (running a R/Julia/Python source file, a bash script, etc..)

**Resources** Resources required to run the job (memory, cpus, time, software requirements...)

The Job/Resource manager allocates the available resources of the computing nodes of the cluster to the different users.



# Running Jobs on a Computing Cluster




**Job** A program to be executed. (running a R/Julia/Python source file, a bash script, etc..)

**Resources** Resources required to run the job (memory, cpus, time, software requirements...)

The Job/Resource manager allocates the available resources of the computing nodes of the cluster to the different users.

# Running Jobs on a Computing Cluster



**Job** A program to be executed. (running a R/Julia/Python source file, a bash script, etc..)

**Resources** Resources required to run the job (memory, cpus, time, software requirements...)

The Job/Resource manager allocates the available resources of the computing nodes of the cluster to the different users.

# Module Loading

Many packages are installed on the cluster. To use them, we use the following special commands.

- Show all modules: `module avail`
- Search for modules: `eo-module-find julia`
- Load a module: `module load engaging/julia/0.6.1`

# Running an Interactive Job on Engaging - `srun`

To ask SLURM, we use the command `srun`.

- The main resource parameters are:

Memory `--mem-per-cpu=1G`

CPUs `--cpus-per-task=1`

Time `--time=1-12:00` (for 1 day and 12 hours)

Partition `--partition=sched_mit_sloan_interactive`

- The partition is the name of the cluster you want to use. The two main partitions you should use are `sched_mit_sloan_interactive` and `sched_mit_sloan_batch`.
- When running interactive jobs, we should use the `--pty` flag to start the session in the current window.

# Monitoring Jobs

- You can check if your job is running or in the queue:  
`eo-show-myjobs`
- Cancel a job: `scancel <JOB_ID>`

# Running a Batch Job on Engaging - sbatch

- A nice way to run jobs is to save them and all their parameters in a special bash file, and use `sbatch` to run them.
- The job can be run using `"sbatch myjobname.sh"`
- The output of your jobs (what would normally appear in the console) will be saved in a `slurm-<yourjob>.out` log file where you ran `sbatch`.

# Types of Parallel Computing Jobs

- “Embarassingly parallel”
  - Multiple instances of the same program with different inputs
  - Each instance can run independently, no need for shared memory
  - eg. grid search for hyperparameter tuning
- Parallel computing with shared memory
  - Commands execute simultaneously on separate processors, but work on shared memory
  - Computation cannot (easily) be split across separate nodes because communication between processes is necessary
  - eg. matrix multiplication, parallel SGD
- Advanced topics
  - Distributed computing with a message passing interface (MPI)
  - GPU computing

## Other Things to Explore

There are a lot of possibilities with Engaging and other systems! Read the Engaging and SLURM documentation, use Google, etc.. For example:

- Running interactive jobs (ie with a console that you control)
- Running graphical interfaces (RStudio, SAS, Matlab...)
- Receive an email when your job completes
- Use Jupyter Notebooks on the cluster...